

TASK_1_Titanic_Survival_Prediction

1. Loading the dataset

```
[1]: import pandas as pd

# GitHub raw URL
url = 'https://raw.githubusercontent.com/abuthahir17/CODSOFT_INTERNSHIP/main/
↳Titanic-Dataset.csv.xlsx'

# Read Excel file
data = pd.read_excel(url)

# Display top rows (First 5 rows)
data.head()
```

```
[1]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

2. Data Inspection

2.1. Checking Data Types

```
[2]: print("Data Types:\n\n" ,data.dtypes)
```

Data Types:

```
PassengerId      int64
Survived          int64
Pclass            int64
Name              object
Sex               object
Age              float64
SibSp             int64
Parch            int64
Ticket            object
Fare              float64
Cabin             object
Embarked          object
dtype: object
```

2.2. Checking Dataset Shape

```
[3]: print("Dataset Shape:", data.shape)
```

Dataset Shape: (891, 12)

2.3. Describing the Dataset

```
[4]: print("Describe the Dataset \n\n" ,data.describe())
```

Describe the Dataset

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

3. Handling Missing Values

3.1. Initial Check for Missing Values

```
[5]: print("\nMissing values:\n\n", data.isnull().sum())
```

Missing values:

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

3.2. Filling Missing 'Embarked' Values

```
[6]: #data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)      #it
      ↪ gives future warning for inplace=True
      data['Embarked'] = data['Embarked'].fillna(data['Embarked'].mode()[0])
```

3.3. Re-checking Missing Values (after Embarked imputation)

```
[7]: print("\nAgain Check Missing values:\n\n", data.isnull().sum())
      #Missing value in Embarked is 0
```

Again Check Missing values:

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        0
dtype: int64
```

3.4 Filling Missing 'Age' Values

```
[8]: # Fill missing age using median by Pclass and Sex group
data['Age'] = data.groupby(['Pclass', 'Sex'])['Age'].transform(lambda x: x.
    ↪ fillna(x.median()))
```

3.5 Re-checking Missing Values (after Age imputation)

```
[9]: print("\nAgain Check Missing values:\n\n", data.isnull().sum())
#Missing value in Age is 0
```

Again Check Missing values:

```
PassengerId      0
Survived          0
Pclass            0
Name              0
Sex               0
Age               0
SibSp             0
Parch             0
Ticket           0
Fare              0
Cabin            687
Embarked          0
dtype: int64
```

3.6 Dropping 'Cabin' Column

```
[10]: # Drop Cabin column because too many missing values
data.drop(columns=['Cabin'], inplace=True)
print("Preprocessed_Data is ready")
```

Preprocessed_Data is ready

4. Saving Preprocessed Data

```
[11]: # Save as Excel
excel_file = 'Preprocessed_data.xlsx'
data.to_excel(excel_file, index=False)
```


5. Data Visualization

5.1. Barplot for Survival by gender

```
[12]: # Plot for Survival by gender
import seaborn as sns
import matplotlib.pyplot as plt

sns.barplot(x='Sex', y='Survived', data=data, color = "Red")
```


```
plt.title("Survival Rate by Gender")
plt.show()
```



5.2. Barplot for Survival by Embarked

```
[13]: # Plot for Survival by Embarked
import seaborn as sns
import matplotlib.pyplot as plt

sns.barplot(x='Embarked', y='Survived', data=data, color = "Purple")
plt.title("Survival Rate by Embarked")
plt.show()
```



5.3. Scatterplot for Age vs Fare (Survival Highlighted)


```
[14]: import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib.lines import Line2D

sns.set_style('whitegrid')

plt.figure(figsize=(7, 5))
sns.scatterplot(x=data['Age'], y=data['Fare'], hue=data['Survived'], palette={0:
    ↪ 'red', 1: 'green'}, alpha=0.7)

# Custom legend with matching colors and text
custom_legend = [
    Line2D([0], [0], marker='o', label='Not Survived', markerfacecolor='red', ↪
    ↪ markersize=8),
    Line2D([0], [0], marker='o', label='Survived', markerfacecolor='green', ↪
    ↪ markersize=8)
]
# Plot titles and labels
```

```
plt.title('Scatterplot of Age vs Fare (Survival Highlighted)')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.legend(handles=custom_legend)
plt.show()
```




5.4. Barplot for Survival by Passenger Class

```
[15]: import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('whitegrid')

# Plot titles and labels
plt.figure(figsize=(7, 5))
sns.barplot(x='Pclass', y='Survived', data=data, color="blue")
plt.title('Survival Rate by Passenger Class')
plt.xlabel('Passenger Class')
plt.ylabel('Survival Rate')
plt.show()
```



6. Feature Engineering

6.1. Dropping Irrelevant Features

```
[16]: # Drop unnecessary columns  
data_model = data.drop(columns=["PassengerId", "Ticket"])
```

6.2. Inspecting DataFrame Columns

```
[17]: data.columns
```

```
[17]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
         'Parch', 'Ticket', 'Fare', 'Embarked'],  
        dtype='object')
```

6.3. Encoding Categorical Variables

```
[18]: from sklearn.preprocessing import LabelEncoder  
  
      # Encode categorical columns  
data_model['Sex'] = LabelEncoder().fit_transform(data_model['Sex'])  
data_model['Embarked'] = LabelEncoder().fit_transform(data_model['Embarked'])
```


7. Model Training and Performance Evaluation

7.1. Data Splitting and Model Training

```
[19]: from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier

      # Features and target
      X = data_model.drop(columns=["Survived", "Name"])
      y = data_model["Survived"]

      # Train-test split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)

      # Train the model
      model = RandomForestClassifier(random_state=42)
      model.fit(X_train, y_train)

      # Predict
      y_pred = model.predict(X_test)
```

7.2. Model Accuracy Assessment

```
[20]: from sklearn.metrics import accuracy_score

      accuracy = accuracy_score(y_test, y_pred)
      print(f"Model Performance: \n\nAccuracy: {accuracy * 100:.2f}%")
      print(f"The model correctly predicts survival {accuracy * 100:.0f} out of 100,
      ↪times.\n\n")
```

Model Performance:

Accuracy: 82.12%

The model correctly predicts survival 82 out of 100 times.

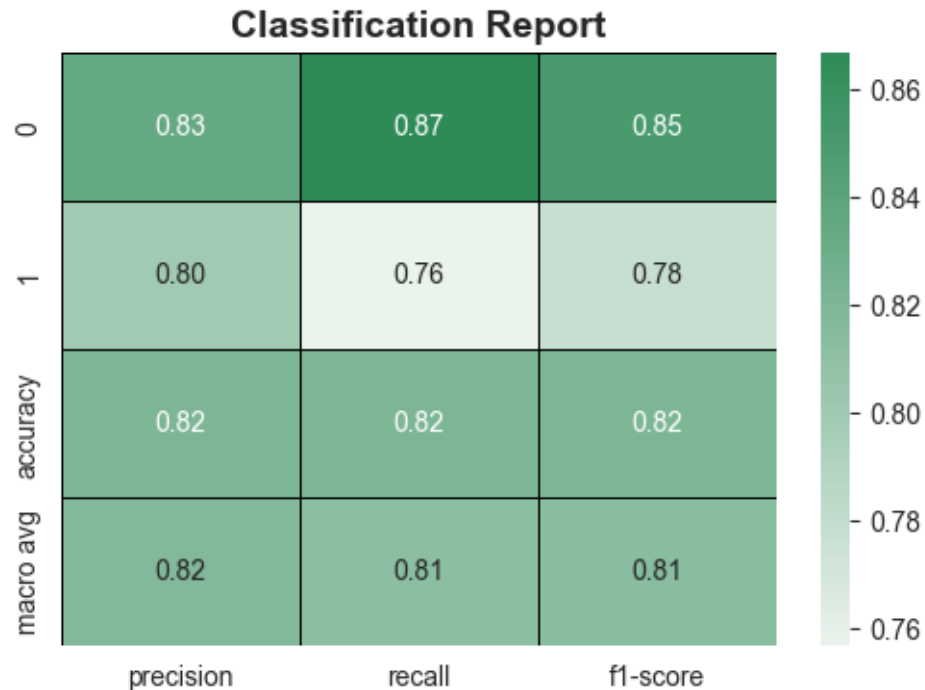
7.3. Classification Report

```
[21]: from sklearn.metrics import classification_report
      import pandas as pd
      import seaborn as sns
      import matplotlib.pyplot as plt

      # Convert classification report to DataFrame
      report_dict = classification_report(y_test, y_pred, output_dict=True)
      report_df = pd.DataFrame(report_dict).transpose()
```

```
plt.figure(figsize=(6, 4))
sns.heatmap(report_df.iloc[: -1, : -1], annot=True, fmt=".2f",
            cmap=sns.light_palette("seagreen", as_cmap=True),
            linewidths=.5, linecolor='black')

plt.title("Classification Report", fontsize=14, weight='bold')
plt.show()
```



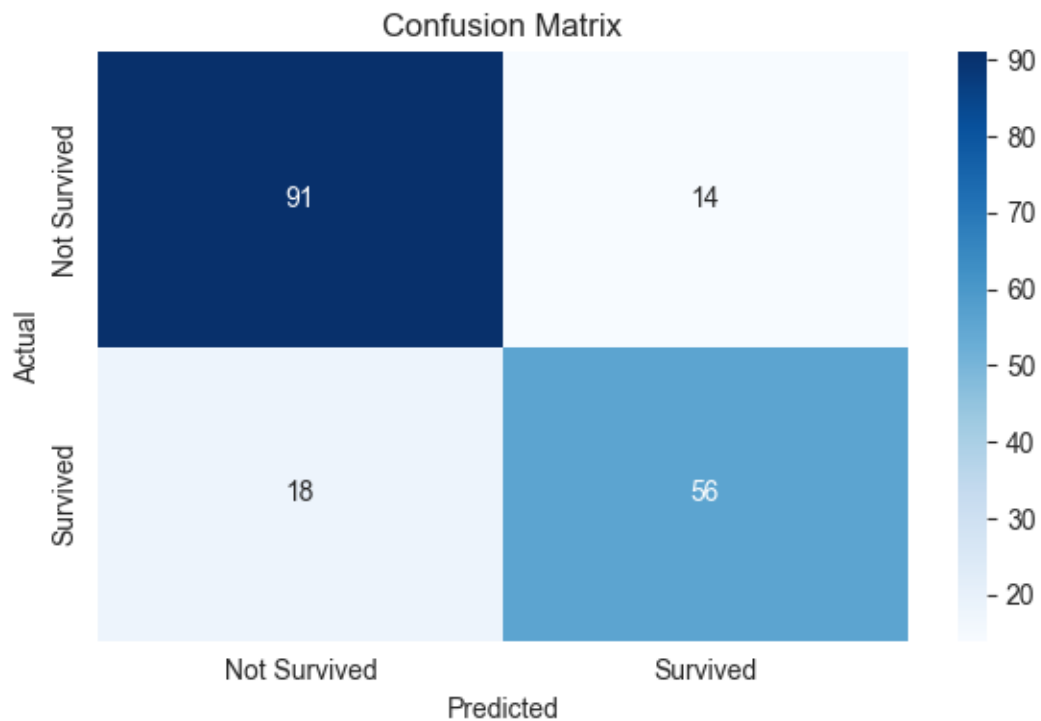
7.4. Confusion Matrix Visualization

```
[22]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Not Survived', 'Survived'],
            yticklabels=['Not Survived', 'Survived'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
```

```
plt.tight_layout()
plt.show()
```




7.5. Feature Correlation Analysis (Heatmap)

```
[23]: import matplotlib.pyplot as plt
import seaborn as sns

# Drop the 'Name' column and any other non-numeric columns not needed for
↳ correlation
# (e.g., 'Ticket' if it exists and is not numeric, though it's less common to
↳ correlate)
data_for_correlation = data_model.drop(columns=['Name', 'Ticket'],
↳ errors='ignore')

correlation_matrix = data_for_correlation.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
↳ linewidths=.5)
plt.title('Correlation Heatmap of Titanic Features')
plt.show()
```



8. Boosting Model Performance (Accuracy)

8.1. Calculating Family Size

```
[24]: data_model['FamilySize'] = data_model['SibSp'] + data_model['Parch'] + 1
```

8.2. Identifying Alone Passengers

```
[25]: #Creating 'IsAlone' based on 'FamilySize'
data_model['IsAlone'] = (data_model['FamilySize'] == 1).astype(int)
```

8.3. Creating Age Bins

```
[26]: #Categorizing 'Age' into 'AgeRange' for better insights
import pandas as pd
data_model['AgeRange'] = pd.cut(data_model['Age'], bins=[0, 12, 18, 30, 45, 60, 100],
```

```
labels=['Child', 'Teenager', 'Young Adult', 'Adult', 'Middle Aged',
        ↪'Senior'])
```

8.4. Extracting and Encoding Titles

```
[27]: # Normalizing 'Name' titles and converting them to numerical format

data_model['Title'] = data_model['Name'].str.extract(r' ([A-Za-z+])\.',
        ↪expand=False)
data_model['Title'] = data_model['Title'].replace(['Mlle', 'Ms'], 'Miss')
data_model['Title'] = data_model['Title'].replace(['Mme', 'Lady'], 'Mrs')
data_model['Title'] = data_model['Title'].replace(['Dr', 'Rev', 'Col', 'Major',
        ↪'Sir', 'Don', 'Jonkheer', 'Capt'], 'Rare')

from sklearn.preprocessing import LabelEncoder
data_model['Title'] = LabelEncoder().fit_transform(data_model['Title'])
```

8.5. Barplot for Survival Rate by Age Group


```
[28]: import matplotlib.pyplot as plt
import pandas as pd

# Step 1: Group total and survived counts by AgeRange
total_by_age = data_model.groupby('AgeRange', observed=True)['Survived'].count()
survived_by_age = data_model.groupby('AgeRange', observed=True)['Survived'].
        ↪sum()

# Calculate survival rate
survival_rate = (survived_by_age / total_by_age) * 100

plt.figure(figsize=(7, 5))
survival_rate.plot(kind='bar', color='skyblue', edgecolor='black')

plt.title("Survival Rate by Age Group")
plt.ylabel("Survival Rate (%)")
plt.xlabel("Age Group")
plt.ylim(0, 100)
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



8.6. Final Model Accuracy Check

[29]: *#Assessing the effectiveness of feature engineering on model performance*

```
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

# Features and target
X = data_model.drop(columns=["Survived", "Name", "AgeRange"])
y = data_model["Survived"]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

# Train the model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Performance: \n\nAccuracy: {accuracy * 100:.2f}%")
print(f"The model correctly predicts survival {accuracy * 100:.0f} out of 100_
↳times.\n\n")

#This will increase the accuracy
```

Model Performance:

Accuracy: 84.36%

The model correctly predicts survival 84 out of 100 times.

```
[30]: print("--- Titanic Survival Prediction Task Complete ---")
```

--- Titanic Survival Prediction Task Complete ---