

TASK_4_SALES_PREDICTION_USING_PYTHON

1. Loading the dataset

```
[1]: import pandas as pd

# GitHub raw URL
url = 'https://raw.githubusercontent.com/abuthahir17/CODSOFT_INTERNSHIP/main/
      ↪advertising.csv'

# Read CSV file
data = pd.read_csv(url)

print("Dataset loaded successfully.")
```

Dataset loaded successfully.

2. Data Inspection

2.1. First 5 Rows

```
[2]: print("First 5 rows of the dataset:")
      print(data.head())
```

First 5 rows of the dataset:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

2.2. Dataset Information

```
[3]: print("Information about the dataset:")
      print(data.info())
```

Information about the dataset:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 200 entries, 0 to 199

Data columns (total 4 columns):

#	Column	Non-Null Count	Dtype
----	-----	-----	-----

```

0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
None

```

2.3. Describing the Dataset

```
[4]: print("Descriptive statistics of the dataset:")
      print(data.describe())
```

```
Descriptive statistics of the dataset:
```

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

2.4. Checking Dataset Shape

```
[5]: print("Dataset Shape:", data.shape)
```

```
Dataset Shape: (200, 4)
```

2.5. Checking Missing Values

```
[6]: print("Missing values in each column:")
      print(data.isnull().sum())

      # There is no missing value
```

```
Missing values in each column:
TV          0
Radio       0
Newspaper   0
Sales       0
dtype: int64
```

3. Data Visualization

3.1. TV Advertising vs. Sales Scatter Plot

```
[7]: import seaborn as sns
      import matplotlib.pyplot as plt

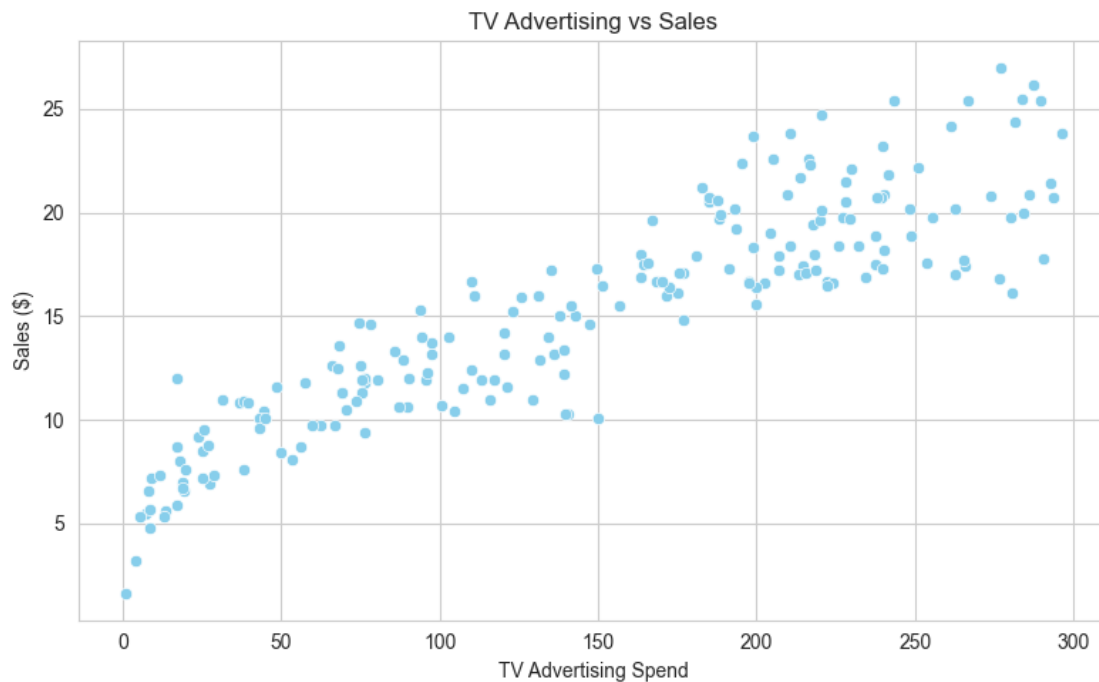
      # Set seaborn style
```

```

sns.set_style("whitegrid")

# TV vs Sales
plt.figure(figsize=(8, 5))
sns.scatterplot(x='TV', y='Sales', data=data, color='skyblue')
plt.title('TV Advertising vs Sales')
plt.xlabel('TV Advertising Spend')
plt.ylabel('Sales ($)')
plt.grid(True)
plt.tight_layout()
plt.show()

```

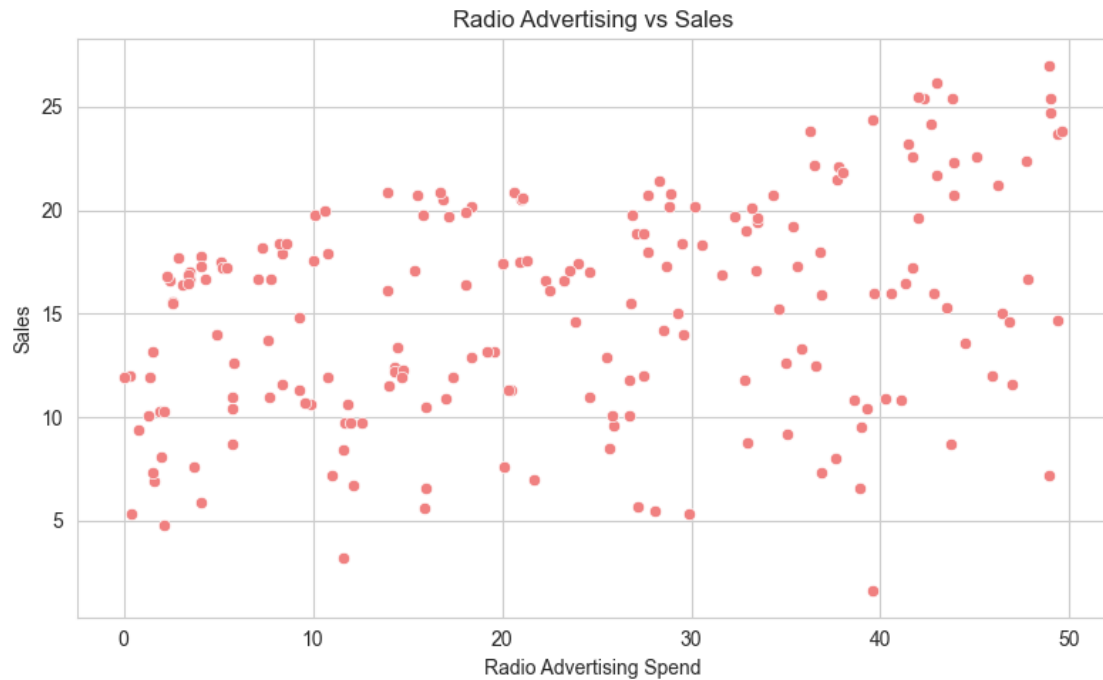


3.2. Radio Advertising vs. Sales Scatter Plot

```

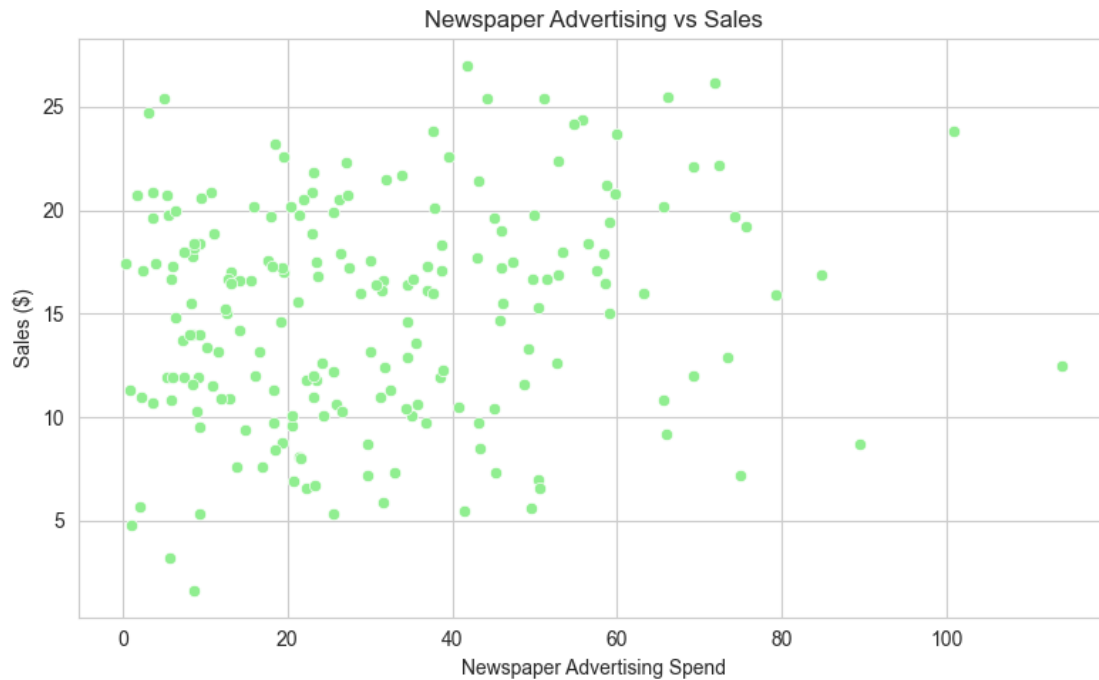
[8]: # Radio vs Sales
plt.figure(figsize=(8,5))
sns.scatterplot(x='Radio', y='Sales', data=data, color='lightcoral')
plt.title('Radio Advertising vs Sales')
plt.xlabel('Radio Advertising Spend')
plt.ylabel('Sales')
plt.grid(True)
plt.tight_layout()
plt.show()

```



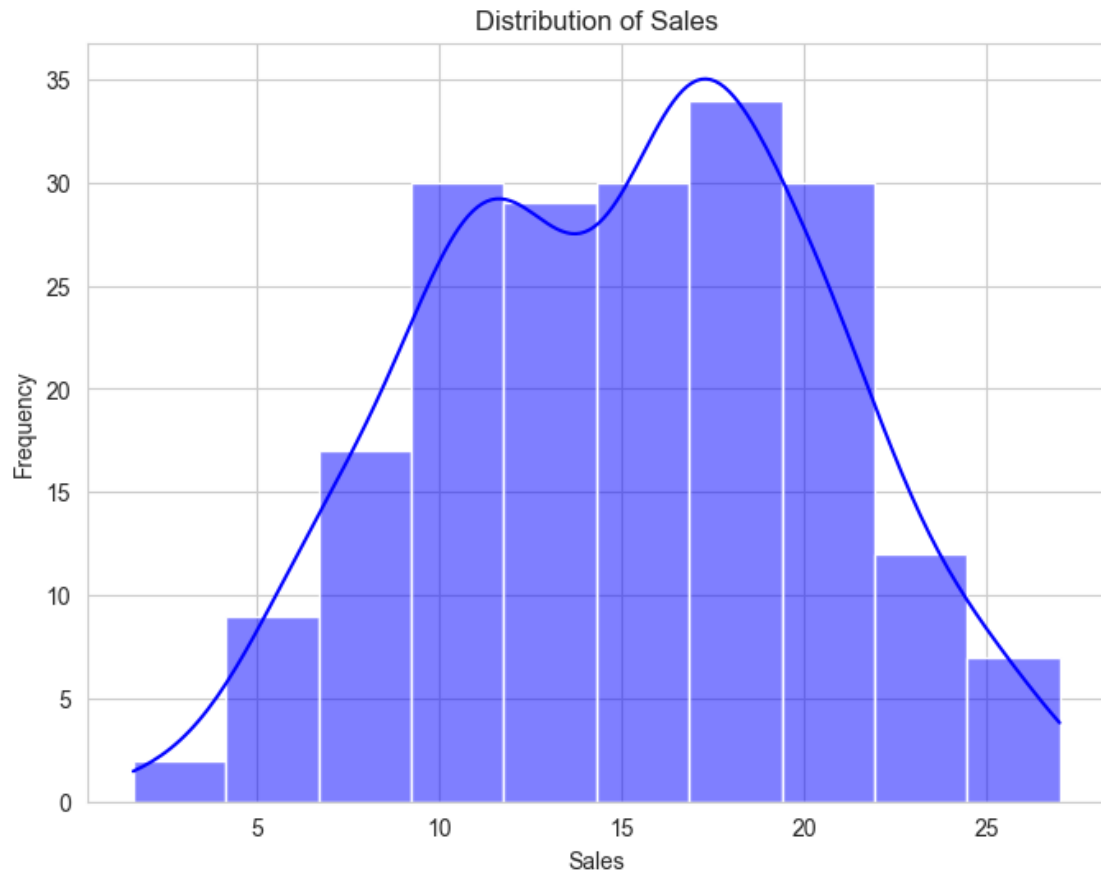
3.3. Newspaper Advertising vs. Sales Scatter Plot

```
[9]: # Newspaper vs Sales
plt.figure(figsize=(8, 5))
sns.scatterplot(x='Newspaper', y='Sales', data=data, color='lightgreen')
plt.title('Newspaper Advertising vs Sales')
plt.xlabel('Newspaper Advertising Spend')
plt.ylabel('Sales ($)')
plt.grid(True)
plt.tight_layout()
plt.show()
```




3.4. Distribution of Sales Histogram

```
[10]: # Distribution of Sales
plt.figure(figsize=(8, 6))
sns.histplot(data['Sales'], kde=True, color='blue')
plt.title('Distribution of Sales')
plt.xlabel('Sales')
plt.ylabel('Frequency')
plt.show()
```



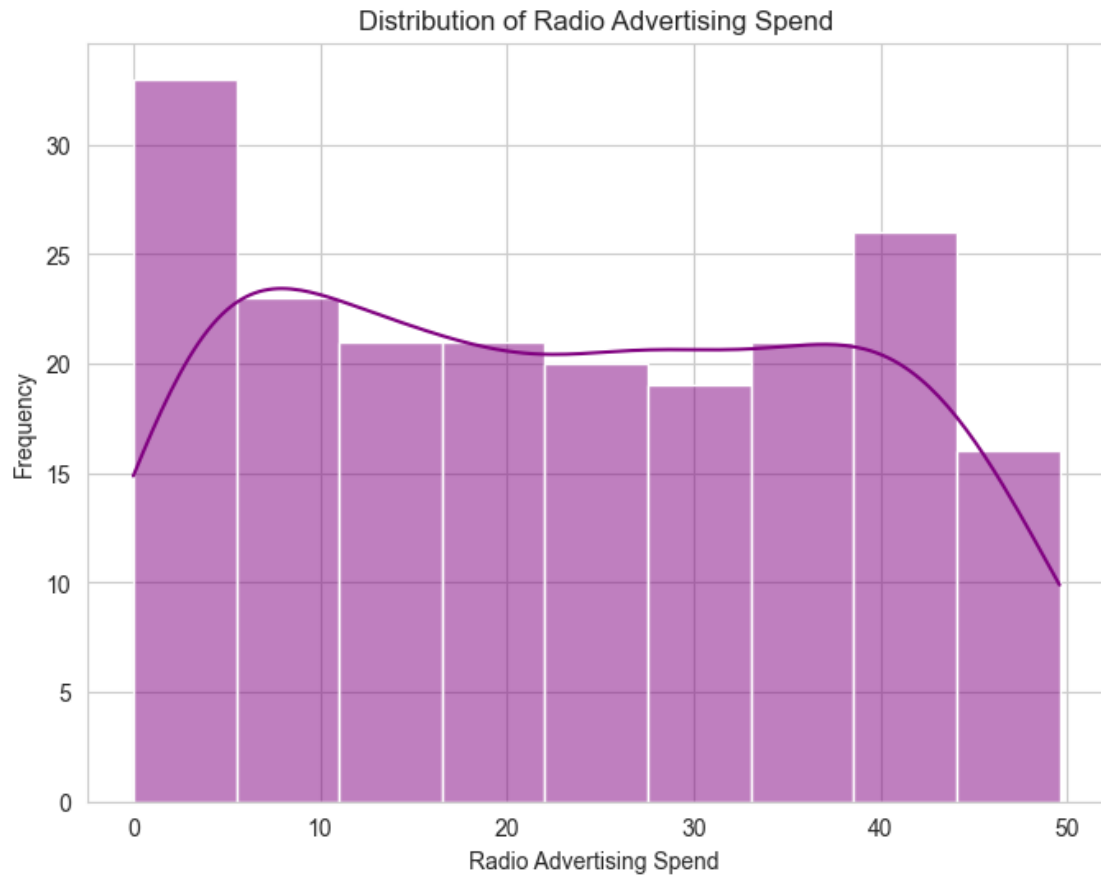
3.5. Distribution of TV Advertising Spend Histogram

```
[11]: # Distribution of TV
plt.figure(figsize=(8, 6))
sns.histplot(data['TV'], kde=True, color = "orange")
plt.title("Distribution of TV Advertising Spend")
plt.xlabel("TV Advertising Spend")
plt.ylabel("Frequency")
plt.show()
```



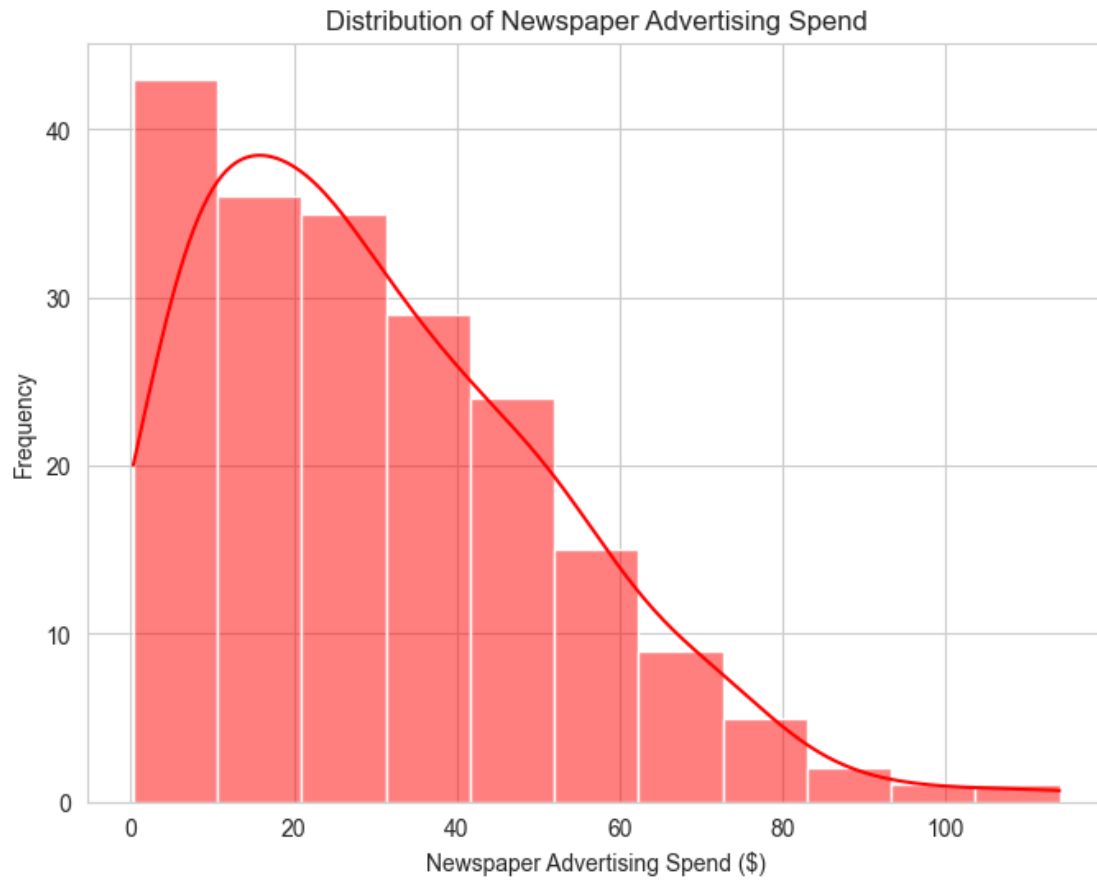
3.6. Distribution of Radio Advertising Spend Histogram

```
[12]: # Distribution of Radio
plt.figure(figsize=(8, 6))
sns.histplot(data['Radio'], kde=True, color = "purple")
plt.title("Distribution of Radio Advertising Spend")
plt.xlabel("Radio Advertising Spend ")
plt.ylabel("Frequency")
plt.show()
```



3.7. Distribution of Newspaper Advertising Spend Histogram

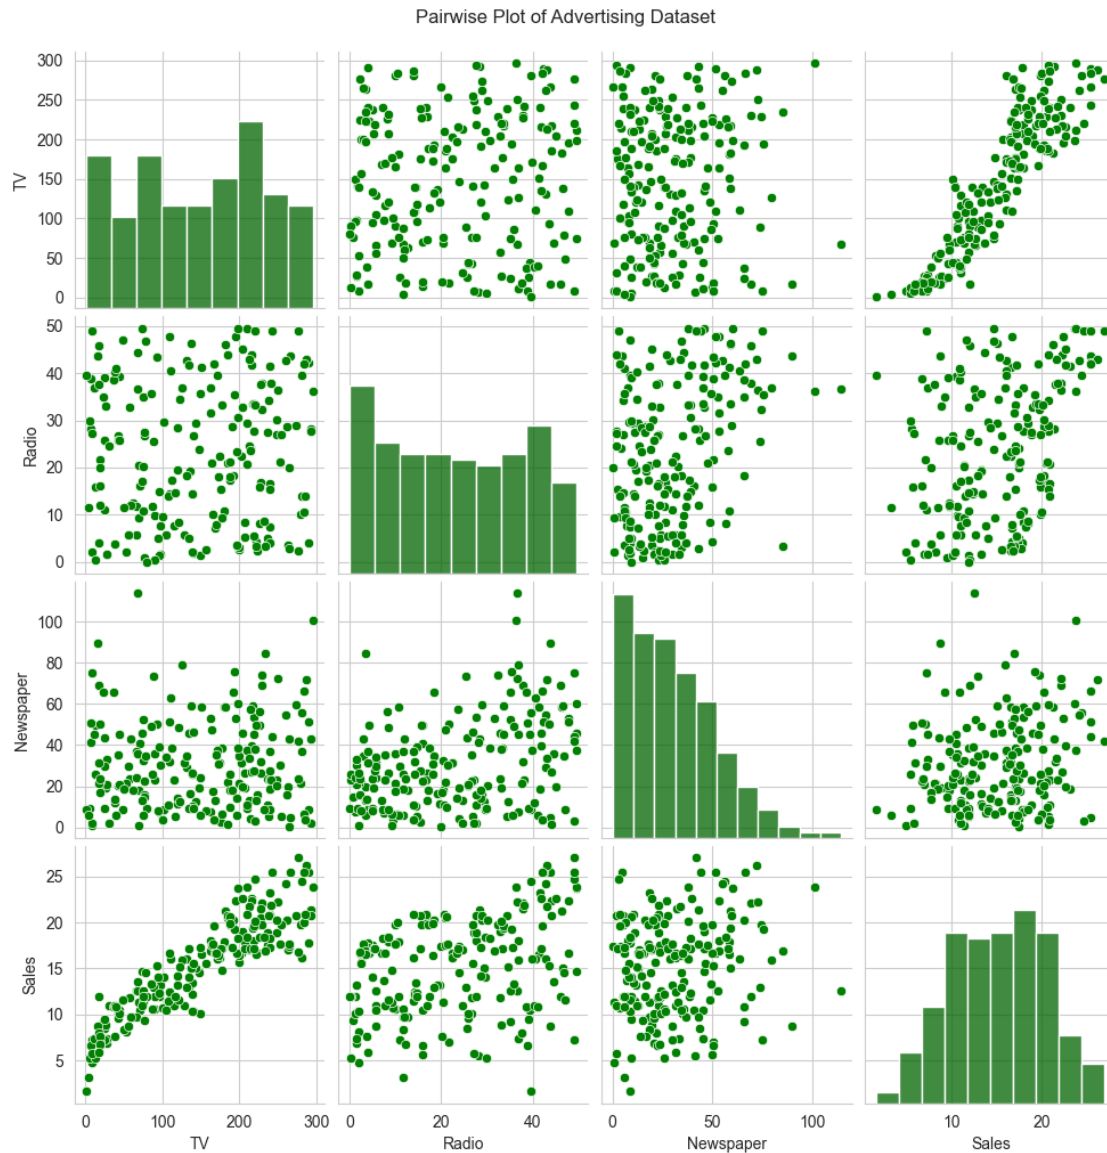
```
[13]: # Distribution of Newspaper
plt.figure(figsize=(8, 6))
sns.histplot(data['Newspaper'], kde=True, color = "red")
plt.title("Distribution of Newspaper Advertising Spend")
plt.xlabel("Newspaper Advertising Spend ($)")
plt.ylabel("Frequency")
plt.show()
```

3.8. Pairwise Plot of Advertising Dataset

```
[14]: import seaborn as sns
import matplotlib.pyplot as plt


sns.pairplot(data, diag_kws={'color': 'darkgreen', 'fill': True},
             plot_kws={'color': 'green'})
plt.suptitle("Pairwise Plot of Advertising Dataset", y=1.02)
plt.show()
```



3.9. Feature Correlation Heatmap

```
[15]: print("Correlation Heatmap:\n")
plt.figure(figsize=(7, 5))
sns.heatmap(data.corr(), annot=True, cmap='YlGnBu')
plt.title("Feature Correlation Heatmap")
plt.show()
```

Correlation Heatmap:



4. Data Preprocessing

4.1. Feature and Target Separation

```
[16]: # Define features (X) and target (y)
# Split the data
X = data[['TV', 'Radio', 'Newspaper']]
y = data['Sales'] # Target variable
```

4.2. Data Splitting

```
[17]: from sklearn.model_selection import train_test_split
# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

print(f"\nTraining set size: {X_train.shape[0]} samples")
print(f"Testing set size: {X_test.shape[0]} samples")
```

Training set size: 160 samples

Testing set size: 40 samples

5. Model Training

5.1. Initialize the models

```
[18]: from sklearn.linear_model import LinearRegression
      # Initialize the Linear Regression model
      # This uses a 'machine learning technique' (Linear Regression) to predict sales.
      model = LinearRegression()

      # Train the model using the training data
      model.fit(X_train, y_train)

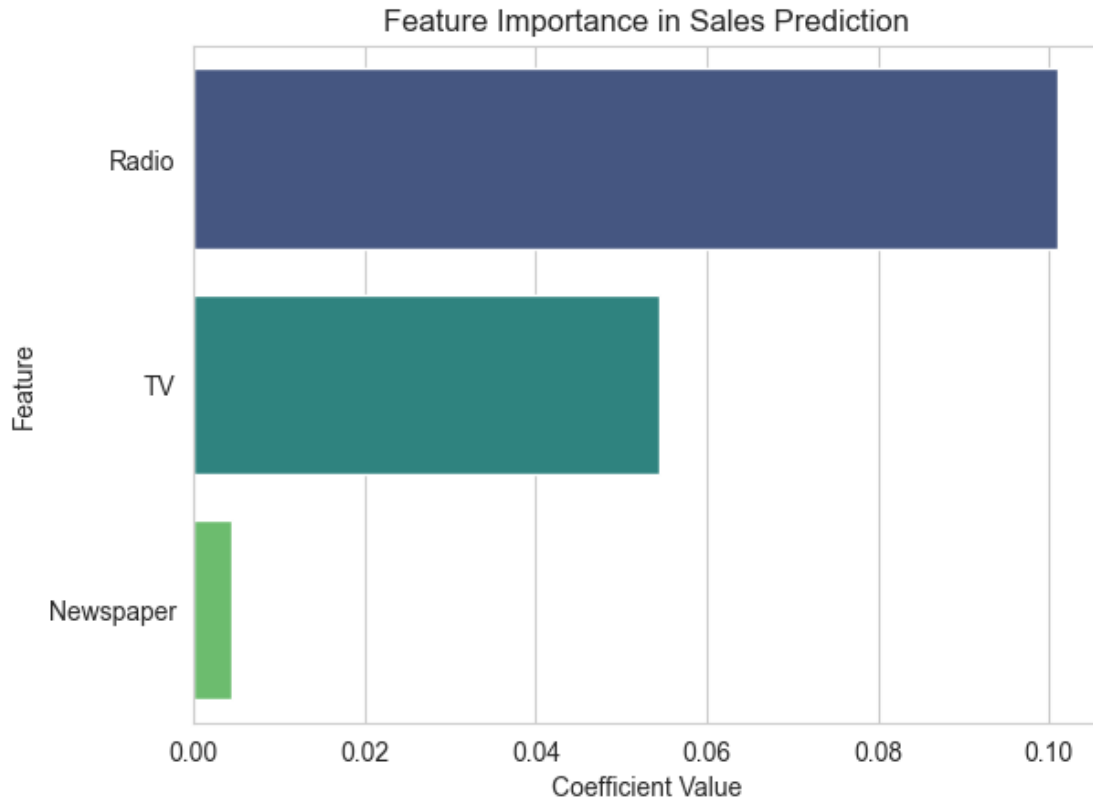
      print("Model training complete.")
```

Model training complete.

5.2. Feature Importance

```
[19]: # Coefficients of the trained Linear Regression model
      coefficients = pd.DataFrame({ 'Feature': X.columns, 'Coefficient': model.coef_
      ↪ }).sort_values(by='Coefficient', ascending=False)

      # Bar chart of coefficients
      sns.barplot(x='Coefficient', y='Feature', data=coefficients, palette='viridis',
      ↪ hue='Feature')
      plt.title("Feature Importance in Sales Prediction")
      plt.xlabel("Coefficient Value")
      plt.ylabel("Feature")
      plt.show()
```



5.3. Model Prediction

```
[20]: y_pred = model.predict(X_test)
```

6. Model Evaluation

6.1. Calculating Regression Model Performance

```
[21]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import numpy as np
# Calculate evaluation metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R-squared (R2): {r2:.2f}")

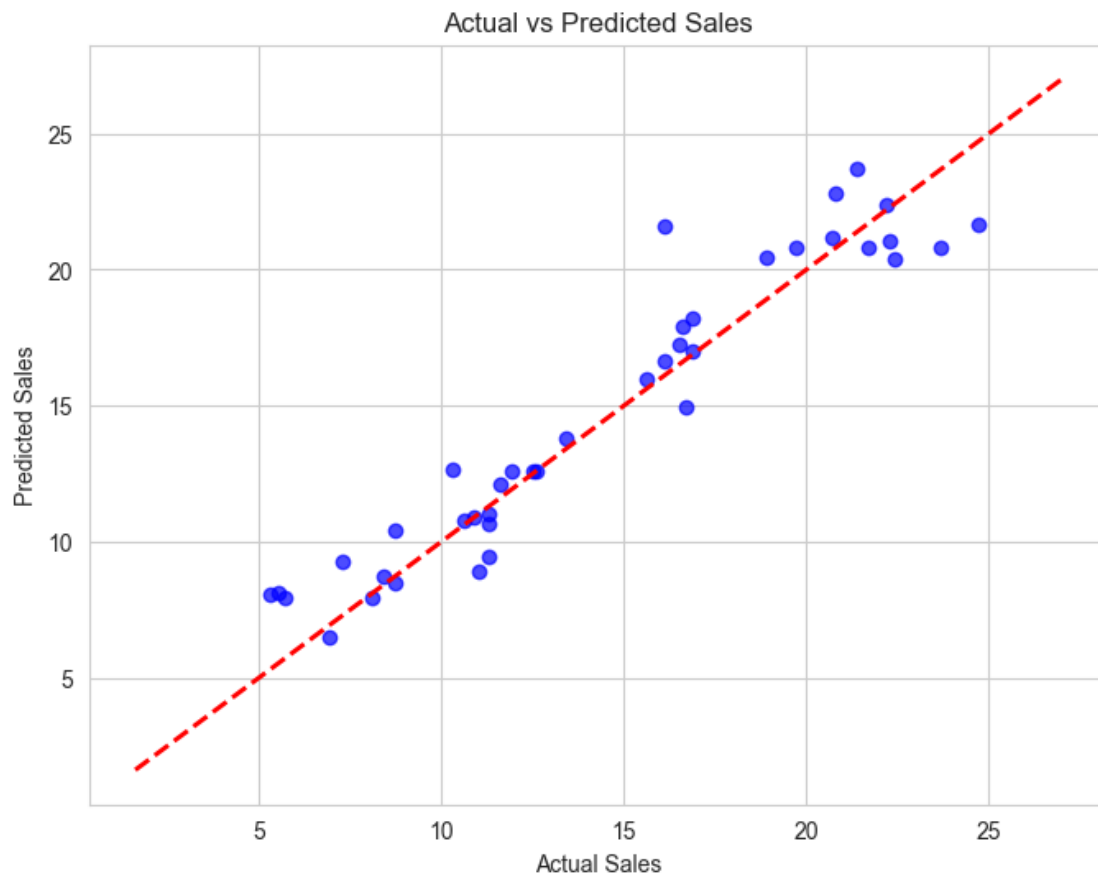
print(f"\nModel Accuracy (based on R-squared): {r2 * 100:.2f}%")
```

Mean Absolute Error (MAE): 1.27
Mean Squared Error (MSE): 2.91
Root Mean Squared Error (RMSE): 1.71
R-squared (R2): 0.91

Model Accuracy (based on R-squared): 90.59%

6.2. Visualizing Actual vs. Predicted Sales

```
[22]: # Plotting actual vs predicted values
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.7, color='blue')
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', lw=2) # Diagonal line
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs Predicted Sales')
plt.show()
```



7. Boosting Model Performance (Accuracy)

7.1. Creating a Binary Sales Target to improve Accuracy

```
[23]: # Convert 'Sales' into a binary classification target: High Sales (1) or Low
      ↪Sales (0)
      data['Sales_Class'] = (data['Sales'] > data['Sales'].median()).astype(int)
```

7.2. Feature and Target Variable Definition

```
[24]: # Features and classification target
      X = data[['TV', 'Radio', 'Newspaper']]
      y = data['Sales_Class']
```

7.3. Splitting Data into Training and Testing Sets

```
[25]: # Split into training and testing
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

7.4. Random Forest Classifier: Training and Evaluation

```
[26]: from sklearn.linear_model import LogisticRegression
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score, confusion_matrix
      # Random Forest Classifier
      rf = RandomForestClassifier(random_state=42)
      rf.fit(X_train, y_train)
      y_pred_rf = rf.predict(X_test)
      acc_rf = accuracy_score(y_test, y_pred_rf)
      cm_rf = confusion_matrix(y_test, y_pred_rf)
```

7.5. Logistic Regression: Training and Evaluation

```
[27]: # Logistic Regression
      logreg = LogisticRegression()
      logreg.fit(X_train, y_train)
      y_pred_log = logreg.predict(X_test)
      acc_log = accuracy_score(y_test, y_pred_log)
      cm_log = confusion_matrix(y_test, y_pred_log)
```

7.6. Displaying Random Forest Performance and Confusion Matrix


```
[28]: print("Random Forest Classifier:")
      print(f"Accuracy: {acc_rf * 100 :.2f}%")

      print("\nConfusion Matrix for Random Forest:\n")
      plt.figure(figsize=(5,4))
      sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Blues', xticklabels=['Low',
      ↪'High'], yticklabels=['Low', 'High'])
      plt.title('Confusion Matrix - Random Forest')
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
      plt.show()
```

Random Forest Classifier:

Accuracy: 97.50%

Confusion Matrix for Random Forest:



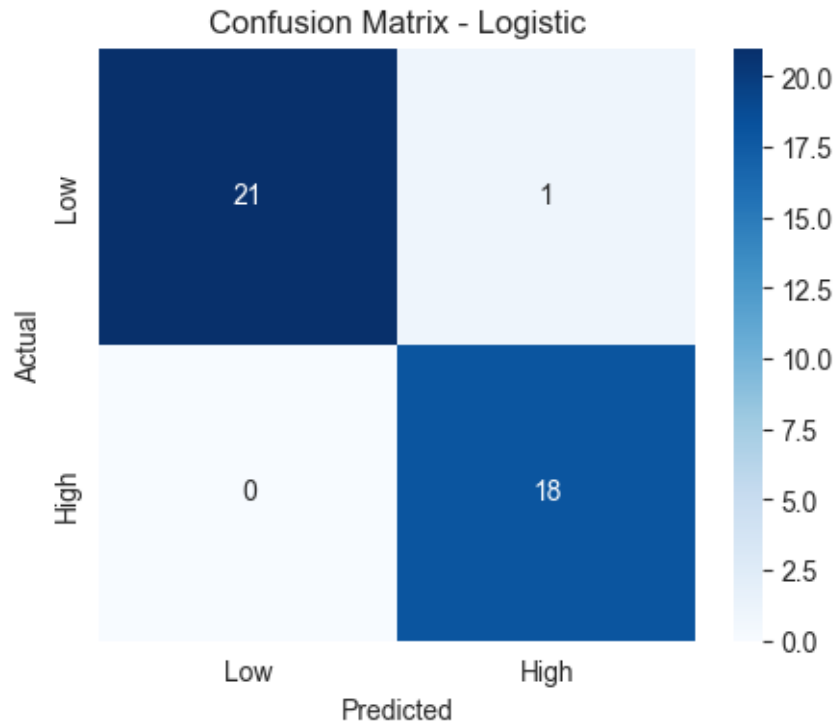
7.7. Displaying Logistic Regression Performance and Confusion Matrix

```
[29]: print("Logistic Regression:")
print(f"Accuracy: {acc_log * 100 :.2f}%")
print("\nConfusion Matrix for Logistic Regression:\n")
plt.figure(figsize=(5,4))
sns.heatmap(cm_log, annot=True, fmt='d', cmap='Blues', xticklabels=['Low', 'High'],
            yticklabels=['Low', 'High'])
plt.title('Confusion Matrix - Logistic')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Logistic Regression:

Accuracy: 97.50%

Confusion Matrix for Logistic Regression:



8. Sales Prediction Example

8.1. Example 1: Forecasting Sales with Specific Advertising Spends

```
[30]: # Example: Predict sales for new advertising expenditures
# Predict sales based on a specific advertising budget

new_ad_spend = pd.DataFrame([[200, 30, 40]], columns=['TV', 'Radio', 'Newspaper'])
predicted_sales = model.predict(new_ad_spend)

print(f"Advertising Spend (TV: 200, Radio: 30, Newspaper: 40):")
print(f"Predicted Sales: {predicted_sales[0]:.2f}")
```

Advertising Spend (TV: 200, Radio: 30, Newspaper: 40):
Predicted Sales: 18.82

8.2. Example 2: Sales Prediction for an Alternative Advertising Scenario

```
[31]: # Another example
# Predict sales based on a specific advertising budget

new_ad_spend_2 = pd.DataFrame([[50, 10, 5]], columns=['TV', 'Radio', 'Newspaper'])
predicted_sales_2 = model.predict(new_ad_spend_2)
```

```
print(f"\nAdvertising Spend (TV: 50, Radio: 10, Newspaper: 5):")
print(f"Predicted Sales: {predicted_sales_2[0]:.2f}")
```

```
Advertising Spend (TV: 50, Radio: 10, Newspaper: 5):
Predicted Sales: 8.47
```

```
[32]: print("--- Sales Prediction Project Complete ---")
```

```
--- Sales Prediction Project Complete ---
```